

# Contents

|  |          |
|--|----------|
| <b>Porting Applications to HIP</b>   | <b>1</b> |
| Hipify Examples . . . . .  | 1        |
| Frontier instructions . . . . .  | 1        |
| Perlmutttr instructions . . . . .  | 1        |
| Exercise 1: Manual code conversion from CUDA to HIP (10 min) . . . . .             | 2        |
| Exercise 2: Code conversion from CUDA to HIP using HIPify tools (10 min) . . . . . | 2        |
| Exercise 3: Mini-App conversion example . . . . .                                  | 3        |
| Makefile option . . . . .  | 4        |

## Porting Applications to HIP

### Hipify Examples

First, get the examples for this lecture

```
git clone git@github.com:olcf/hip-training-series.git
```

#### Frontier instructions

For the first interactive example, get an slurm interactive session on Frontier (see further below for NERSC Perlmutttr):

```
salloc -N 1 -p batch --reservation=hip_training_2023_8_28 --gpus=1 -t 10:00 -A <project>
```

Outside the reservation window or if you're not on the reservation list, you can do: `salloc -N 1 -p batch --gpus=1 -t 10:00 -A <project>`

Use your project id in the project field. If you do not remember it, run the command without the -A option and it should report your valid projects.

The environment needs to be set up for the rocm software such as hipify-perl and hipcc. Here are the commands for Frontier.

```
module load PrgEnv-amd
module load amd
module load cmake
```

#### Perlmutttr instructions

During the training session node reservation hours, get a slurm interactive session with

```
salloc -N 1 -q shared -C gpu -c 32 -G 1 -t 30:00 -A ntrain8 --reservation=hip_aug28
```

Outside the reservation hours, use

```
salloc -N 1 -q interactive -C gpu -c 32 -G 1 -t 30:00 -A <a project>
```

use your own project instead of ntrain8 if you have a NERSC regular project.

The modules needed for Perlmutttr are slightly different than Frontier. Use these instead. We also need to add the path to the bin directory.

```
module load PrgEnv-gnu/8.3.3
module load hip/5.4.3
module load PrgEnv-nvidia/8.3.3
module load cmake
```

```
export PATH=${PATH}:${HIP_PATH}
```

### Exercise 1: Manual code conversion from CUDA to HIP (10 min)

Choose one or more of the CUDA samples in `hip-training-series/Lecture2/HIPIFY/mini-nbody/cuda` directory. Manually convert it to HIP. Tip: for example, the `cudaMalloc` will be called `hipMalloc`. You can choose from `nbody-block.cu`, `nbody-orig.cu`, or `nbody-soa.cu`

You'll want to compile on the node you've been allocated so that `hipcc` will choose the correct GPU architecture.

### Exercise 2: Code conversion from CUDA to HIP using HIPify tools (10 min)

Use the `hipify-perl` script to "hipify" the CUDA samples you used to manually convert to HIP in Exercise 1. `hipify-perl` is in `$ROCM_PATH/hip/bin` directory and should be in your path.

First test the conversion to see what will be converted

```
hipify-perl -examine nbody-orig.cu
```

You'll see the statistics of HIP APIs that will be generated. The output might be different depending on the ROCm version.

```
[HIPIFY] info: file 'nbody-orig.cu' statistics:
  CONVERTED refs count: 7
  TOTAL lines of code: 91
  WARNINGS: 0
[HIPIFY] info: CONVERTED refs by names:
  cudaFree => hipFree: 1
  cudaMalloc => hipMalloc: 1
  cudaMemcpyDeviceToHost => hipMemcpyDeviceToHost: 1
  cudaMemcpyHostToDevice => hipMemcpyHostToDevice: 1
```

`hipify-perl` is in `$ROCM_PATH/hip/bin` directory and should be in your path. In some versions of ROCm, the script is called `hipify-perl.sh`. Perlmutter does not currently have the `hipify` scripts in the `$ROCM_PATH/hip/bin` directory. We have included them in the `hip-training-series/Lecture2/HIPIFY` directory for these exercises.

Now let's actually do the conversion.

```
hipify-perl nbody-orig.cu > nbody-orig.cpp
```

Compile the HIP programs.

```
hipcc -DSHMOO -I ../ nbody-orig.cpp -o nbody-orig
```

The `#define SHMOO` fixes some timer printouts. Add `--offload-arch=<gpu_type>` to specify the GPU type and avoid the autodetection issues when running on a single GPU on a node.

- Fix any compiler issues, for example, if there was something that didn't hipify correctly.
- Be on the lookout for hard-coded Nvidia specific things like warp sizes and PTX.

Run the program

```
srun ./nbody-orig
```

A batch version of Exercise 2 for Frontier is given below. The batch scripts are also located in the `mini-nbody` directory. Please check them and modify them for your project and the reservation if there is one.

```
#!/bin/bash
#SBATCH -N 1
#SBATCH --ntasks=1
#SBATCH --gpus=1
#SBATCH -p batch
#SBATCH -t 00:10:00
```

```

#SBATCH -A <project id>
#SBATCH --reservation hip_training_2023_8_28

module load PrgEnv-amd
module load amd
module load cmake

cd $HOME/hip-training-series/Lecture2/HIPIFY/mini-nbody/cuda
hipify-perl -print-stats nbody-orig.cu > nbody-orig.cpp
hipcc -DSHMOO -I ../ nbody-orig.cpp -o nbody-orig
srun ./nbody-orig
cd ../../..

```

For Perlmutter, there are a few differences in the batch arguments as well as the modules to be loaded.

```

#!/bin/bash
#SBATCH -N 1
#SBATCH -q shared
#SBATCH -C gpu
#SBATCH -c 32
#SBATCH -G 1
#SBATCH -t 00:30:00
#SBATCH -A ntrain8
#SBATCH --reservation hip_aug28

module load PrgEnv-gnu/8.3.3
module load hip/5.4.3
module load PrgEnv-nvidia/8.3.3
module load cmake

export PATH=${PATH}:${HIP_PATH}

cd $HOME/hip-training-series/Lecture2/HIPIFY/mini-nbody/cuda

../../hipify-perl -print-stats nbody-orig.cu > nbody-orig.cpp
hipcc -DSHMOO -I ../ nbody-orig.cpp -o nbody-orig
srun ./nbody-orig
rm nbody-orig nbody-orig.cpp

```

Notes:

- Hipify tools do not check correctness
- `hipconvertinplace-perl` is a convenience script that does `hipify-perl -inplace -print-stats` command

### Exercise 3: Mini-App conversion example

Load the proper environment. For Frontier, the 5.5.1 version is needed to get the right interfaces.

```

module load PrgEnv-amd
module load amd/5.5.1
module load cmake
export CXX=${ROCM_PATH}/llvm/bin/clang++

```

For Perlmutter, the following environment should be set up.

```

module load PrgEnv-gnu/8.3.3
module load hip/5.4.3

```

```
module load PrgEnv-nvidia/8.3.3
module load cmake
export PATH=${PATH}:${HIP_PATH}
```

The original CUDA version of Pennant has been downloaded and installed at ~/hip-training-series/Lecture2/HIPIFY/Pennant-orig. Only one of the test problems has been included to save disk space. The original source is at

<https://asc.llnl.gov/sites/asc/files/2020-09/pennant-singlenode-cude.tgz>

```
cd ~/hip-training-series/Lecture2/HIPIFY/Pennant-orig
```

```
./hipexamine-perl.sh
```

And review the output

Now do the actual conversion. We want to do the conversion for the whole directory tree, so we'll use hipconvertinplace-sh

```
./hipconvertinplace-perl.sh
```

We want to use .hip extensions rather than .cu, so change all files with .cu to .hip

```
mv src/HydroGPU.cu src/HydroGPU.hip
```

Now we have two options to convert the build system to work with both ROCm and CUDA

## Makefile option

First cut at converting the Makefile. Testing with make can help identify the next step.

- Change CUDACFLAGS to HIPCFLAGS sed -i -e 's/CUDACFLAGS/HIPCFLAGS/g' Makefile
- Change all occurrences of CUDA to HIP sed -i -e 's/CUDA/HIP/g' Makefile
- Change the CXX variable icpc to clang++ located in \${ROCM\_PATH}/llvm/bin/clang++ sed -i -e '/CXX/s/icpc/amdclang++/' Makefile
- Change HIPCC from nvcc to hipcc sed -i -e 's/nvcc/hipcc/' Makefile
- Remove -fast and -fno-alias sed -i -e 's/-fast -fno-alias//' Makefile
- Change all %.cu to %.hip in the Makefile sed -i -e 's/%.cu/%.hip/g' Makefile
- Remove -arch=sm\_21 -ptxas-options=-v sed -i -e 's/-arch=sm\_21 --ptxas-options=-v//' Makefile
- Change the LDFLAGS to those needed for AMD sed -i -e 's/^LDFLAGS/LDFLAGS\_CUDA/' Makefile  
sed -i -e '/^LDFLAGS\_CUDA/aLDFLAGS := -L\${ROCM\_PATH}/hip/lib -lamdhip64' Makefile

This new makefile has a separate compile path for .hip and .cpp files. The .hip files are compiled with hipcc and the .cpp files are compiled with amdclang++. There are different strategies that can be used. A simpler approach is to just compile everything with hipcc. There are some limitations with using hipcc everywhere. One of them is that hipcc cannot be used to compile OpenMP code in a .cpp file. That may limit hybrid GPU programming approaches.

The resulting makefile is included as Makefile.twopath and the all hipcc approach is in Makefile.allhipcc. The makefile supports this just by defining CXX to hipcc.

Now we are just getting compile errors from the source files. We will have to do fixes there. We'll tackle them one-by-one.

The first errors are related to the double2 type.

```
compiling src/HydroGPU.hip
```

```
(CPATH=;hipcc -O3 -I. -c -o build/HydroGPU.o src/HydroGPU.hip)
```

```
In file included from src/HydroGPU.hip:14:
```

```
In file included from src/HydroGPU.hh:16:
```

```
src/Vec2.hh:35:8: error: definition of type 'double2' conflicts with type alias of the same name
```

```
struct double2
```

```
^  
/opt/rocm-5.6.0/include/hip/amd_detail/amd_hip_vector_types.h:1098:1: note: 'double2' declared here  
__MAKE_VECTOR_TYPE__(double, double);  
^
```

```
/opt/rocm-5.6.0/include/hip/amd_detail/amd_hip_vector_types.h:1062:15: note: expanded from macro '__MAKE_VECTOR_TYPE__'  
    using CUDA_name##2 = HIP_vector_type<T, 2>;\br/>    ^
```

```
<scratch space>:316:1: note: expanded from here  
double2
```

HIP defines double2. Let's look at Vec2.hh. At line 33 where the first error occurs. We see an `#ifndef __CUDAACC__` around a block of code there. We also need the `#ifdef` to include HIP as well. Let's check the available compiler defines from the presentation to see what is available. It looks like we can use `__HIP_DEVICE_COMPILE__` or maybe `__HIPCC__`.

- Change line 33 in Vec2.hh to `#ifndef __HIPCC__`

```
sed -e -i 's/#ifndef __CUDAACC__/#ifndef __HIPCC__/' src/Vec2.hh
```

The next error is about function attributes that are incorrect for device code.

```
compiling src/HydroGPU.hip
```

```
(CPATH=;hipcc -O3 -I. -c -o build/HydroGPU.o src/HydroGPU.hip  
src/HydroGPU.hip:168:23: error: no matching function for call to 'cross'  
    double sa = 0.5 * cross(px[p2] - px[p1], zx[z] - px[p1]);  
                        ^~~~~
```

```
src/Vec2.hh:206:15: note: candidate function not viable: call to __host__ function from __device__ function
```

The `FNQUALIFIER` macro is what handles the attributes in the code. We find that defined at line 22 and again we see a `#ifdef __CUDAACC__`. It is another `ifdef` for `__CUDAACC__`. We can see that we need to pay attention to all the CUDA `ifdef` statements.

- Change line 22 to `#ifdef __HIPCC__`

```
sed -e -i 's/#ifdef __CUDAACC__/#ifdef __HIPCC__/' src/Vec2.hh
```

Finally we get an error about already defined operators on double2 types. These appear to be defined in HIP, but not in CUDA. So we change line 84

```
compiling src/HydroGPU.hip
```

```
(CPATH=;hipcc -O3 -I. -c -o build/HydroGPU.o src/HydroGPU.hip)  
src/HydroGPU.hip:149:15: error: use of overloaded operator '+' is ambiguous (with operand types 'double2' and 'double2')  
    zxtot += ctemp2[sn];  
    ~~~~~ ^ ~~~~~
```

```
/opt/rocm-5.6.0/include/hip/amd_detail/amd_hip_vector_types.h:510:26: note: candidate function  
    HIP_vector_type& operator+=(const HIP_vector_type& x) noexcept  
    ^
```

```
src/Vec2.hh:88:17: note: candidate function  
inline double2& operator+=(double2& v, const double2& v2)
```

- Change line 85 to `#elif defined(__CUDAACC__)`

```
sed -i -e '85,85s/#else/#elif defined(__CUDAACC__)/' src/Vec2.hh
```

Now we start getting errors for HydroGPU.hip. The first is for the `atomicMin` function. It is already defined in HIP, so we need to add an `ifdef` for CUDA around the code.

```
compiling src/HydroGPU.hip
```

```
(CPATH=;hipcc -O3 -I. -c -o build/HydroGPU.o src/HydroGPU.hip)  
src/HydroGPU.hip:725:26: error: static declaration of 'atomicMin' follows non-static declaration
```

```
static __device__ double atomicMin(double* address, double val)
~
/opt/rocm-5.6.0/include/hip/amd_detail/amd_hip_atomic.h:478:8: note: previous definition is here
double atomicMin(double* addr, double val) {
~
```

1 error generated when compiling for gfx90a.

- Add `#ifndef __CUDAACC__`/`endif` to the block of code in `HydroGPU.hip` from line 725 to 737

```
sed -i -e '724,724a#ifndef __CUDAACC__' -e '738,738a#endif' src/HydroGPU.hip
```

We finally got through the compiler errors and move on to link errors

linking build/pennant

```
/opt/rocm-5.6.0/llvm/bin/clang++ -o build/pennant build/ExportGold.o build/ImportGMV.o build/Parallel.o
ld.lld: error: undefined symbol: hydroInit(int, int, int, int, int, double, double, double, double, double)
>>> referenced by Hydro.cc
>>> build/Hydro.o:(Hydro::Hydro(InputFile const*, Mesh*))
```

```
ld.lld: error: undefined symbol: hydroGetData(int, int, double2*, double*, double*, double*)
>>> referenced by Hydro.cc
>>> build/Hydro.o:(Hydro::getData())
```

This one is a little harder. We can get more information by using `nm build/Hydro.o |grep hydroGetData` and `nm build/HydroGPU.o |grep hydroGetData`. We can see that the subroutine signatures are slightly different due to the `double2` type on the host and GPU. You can also switch the compiler from `clang++` to `g++` to get a slightly more informative error. We are in a tough spot here because we need the `hipmemcpy` in the body of the subroutine, but the types for `double2` are for the device instead of the host. One solution is to just compile and link everything with `hipcc`, but we really don't want to do that if only one routine needs to use the device compiler. So we cheat by declaring the prototype arguments as `void *` and casting the type in the call with `(void *)`. The types are really the same and it is just arguing with the compiler.

```
nm build/Hydro.o |grep hydroGetData
U _Z12hydroGetDataiiP7double2PdS1_S1_
nm build/HydroGPU.o |grep hydroGetData
0000000000003750 T _Z12hydroGetDataiiP15HIP_vector_typeIdLj2EEPdS2_S2_
```

In `HydroGPU.hh`

- Change line 38 and 39 to from `const double2*` to `const void*` `sed -i -e '38,39s/const double2/const void/' src/HydroGPU.hh`
- Change line 62 from `double2*` to `void*` `sed -i -e '62,62s/double2/void/' src/HydroGPU.hh`

In `HydroGPU.hip`

- Change line 1031 and 1032 to `const void*` `sed -i -e '1031,1032s/const double2/const void/' src/HydroGPU.hip`
- Change line 1284 to `const void*` `sed -i -e '1284,1284s/double2/void/' src/HydroGPU.hip`

In `Hydro.cc`

- Add `(void *)` before the arguments on lines 59, 60, and 145 `sed -i -e '59,59s/mesh/(void *)mesh/' src/Hydro.cc` `sed -i -e '60,60s/pu/(void *)pu/' src/Hydro.cc` `sed -i -e '145,145s/mesh/(void *)mesh/' src/Hydro.cc`

Now it compiles and we can test the run with

```
build/pennant test/sedovbig/sedovbig.pnt
```

So we have the code converted to HIP and fixed the build system for it. But we haven't accomplished our original goal of running with both ROCm and CUDA.

We can copy a sample portable Makefile from `hip-training-series/Lecture1/HIP/saxpy/Makefile` and modify it for this application.

```
EXECUTABLE = pennant
BUILDDIR := build
SRCDIR = src
all: $(BUILDDIR)/$(EXECUTABLE)

.PHONY: test

OBJECTS = $(BUILDDIR)/Driver.o $(BUILDDIR)/GenMesh.o $(BUILDDIR)/HydroBC.o
OBJECTS += $(BUILDDIR)/ImportGMV.o $(BUILDDIR)/Mesh.o $(BUILDDIR)/PolyGas.o
OBJECTS += $(BUILDDIR)/TTS.o $(BUILDDIR)/main.o $(BUILDDIR)/ExportGold.o
OBJECTS += $(BUILDDIR)/Hydro.o $(BUILDDIR)/HydroGPU.o $(BUILDDIR)/InputFile.o
OBJECTS += $(BUILDDIR)/Parallel.o $(BUILDDIR)/QCS.o $(BUILDDIR)/WriteXY.o

CXXFLAGS = -g -O3 -DNDEBUG -fPIC
HIPCC_FLAGS = -O3 -g -DNDEBUG

HIP_PLATFORM ?= amd

ifeq ($(HIP_PLATFORM), nvidia)
    HIP_PATH ?= $(shell hipconfig --path)
    HIPCC_FLAGS += -x cu -I${HIP_PATH}/include/
endif
ifeq ($(HIP_PLATFORM), amd)
    HIPCC_FLAGS += -x hip -munsafe-fp-atomics
endif

$(BUILDDIR)/%.d : $(SRCDIR)/%.cc
    @echo making depends for $<
    $(maketargetdir)
    $(CXX) $(CXXFLAGS) $(CXXINCLUDES) -M $< | sed "1s![^ \t]\+\.o!$(@:.d=.o) $@!" >$@

$(BUILDDIR)/%.d : $(SRCDIR)/%.hip
    @echo making depends for $<
    $(maketargetdir)
    @hipcc $(HIPCCFLAGS) $(HIPCCINCLUDES) -M $< | sed "1s![^ \t]\+\.o!$(@:.d=.o) $@!" >$@

$(BUILDDIR)/%.o : $(SRCDIR)/%.cc
    @echo compiling $<
    $(maketargetdir)
    $(CXX) $(CXXFLAGS) $(CXXINCLUDES) -c -o $@ $<

$(BUILDDIR)/%.o : $(SRCDIR)/%.hip
    @echo compiling $<
    $(maketargetdir)
    hipcc $(HIPCC_FLAGS) -c $^ -o $@

$(BUILDDIR)/$(EXECUTABLE) : $(OBJECTS)
    @echo linking $@
    $(maketargetdir)
    hipcc $(OBJECTS) $(LDFLAGS) -o $@
```

```
test : $(BUILDDIR)/$(EXECUTABLE)
      $(BUILDDIR)/$(EXECUTABLE) test/sedovbig/sedovbig.pnt
```

```
define maketargetdir
  -@mkdir -p $(dir $@) > /dev/null 2>&1
endef
```

```
clean :
  rm -rf $(BUILDDIR)
```

To test the makefile,

```
make
make test
```

To test the makefile build system with CUDA. The pennant run on Frontier does seem to hang at the end of the run. The instructions for Frontier are included in a file at `~/hip-training-series/Lecture2/HIPIFY/frontier_pennant_set`

```
module load cuda
HIP_PLATFORM=nvidia CXX=g++ make
```

To create a cmake build system, we can copy a sample portable Makefile from `hip-training-series/HIP/saxpy/CMakeLists.txt` and modify it for this application.

```
cmake_minimum_required(VERSION 3.21 FATAL_ERROR)
project(Pennant LANGUAGES CXX)
include(CTest)
```

```
set (CMAKE_CXX_STANDARD 14)
```

```
if (NOT CMAKE_BUILD_TYPE)
  set(CMAKE_BUILD_TYPE RelWithDebInfo)
endif(NOT CMAKE_BUILD_TYPE)
```

```
string(REPLACE -O2 -O3 CMAKE_CXX_FLAGS_RELWITHDEBINFO ${CMAKE_CXX_FLAGS_RELWITHDEBINFO})
```

```
if (NOT CMAKE_GPU_RUNTIME)
  set(GPU_RUNTIME "ROCM" CACHE STRING "Switches between ROCM and CUDA")
else (NOT CMAKE_GPU_RUNTIME)
  set(GPU_RUNTIME "${CMAKE_GPU_RUNTIME}" CACHE STRING "Switches between ROCM and CUDA")
endif (NOT CMAKE_GPU_RUNTIME)
# Really should only be ROCM or CUDA, but allowing HIP because it is the currently built-in option
set(GPU_RUNTIMES "ROCM" "CUDA" "HIP")
if(NOT "${GPU_RUNTIME}" IN_LIST GPU_RUNTIMES)
  set(ERROR_MESSAGE "GPU_RUNTIME is set to \"${GPU_RUNTIME}\".\nGPU_RUNTIME must be either HIP, ROCM,
  message(FATAL_ERROR ${ERROR_MESSAGE})
endif()
# GPU_RUNTIME for AMD GPUs should really be ROCM, if selecting AMD GPUs
# so manually resetting to HIP if ROCM is selected
if (${GPU_RUNTIME} MATCHES "ROCM")
  set(GPU_RUNTIME "HIP")
endif (${GPU_RUNTIME} MATCHES "ROCM")
set_property(CACHE GPU_RUNTIME PROPERTY STRINGS ${GPU_RUNTIMES})
```

```
enable_language(${GPU_RUNTIME})
set(CMAKE_${GPU_RUNTIME}_EXTENSIONS OFF)
set(CMAKE_${GPU_RUNTIME}_STANDARD_REQUIRED ON)
```



```

set(PENNANT_CXX_SRCS src/Driver.cc src/ExportGold.cc src/GenMesh.cc src/Hydro.cc src/HydroBC.cc
                    src/ImportGMV.cc src/InputFile.cc src/Mesh.cc src/Parallel.cc src/PolyGas.cc
                    src/QCS.cc src/TTS.cc src/WriteXY.cc src/main.cc)

set(PENNANT_HIP_SRCS src/HydroGPU.hip)

add_executable(pennant ${PENNANT_CXX_SRCS} ${PENNANT_HIP_SRCS} )

# Make example runnable using ctest
add_test(NAME Pennant COMMAND pennant ../test/sedovbig/sedovbig.pnt )
set_property(TEST Pennant PROPERTY PASS_REGULAR_EXPRESSION "End cycle    3800, time = 9.64621e-01")

set(ROCMCC_FLAGS "${ROCMCC_FLAGS} -munsafe-fp-atomics")
set(CUDACC_FLAGS "${CUDACC_FLAGS} ")

if (${GPU_RUNTIME} MATCHES "HIP")
    set(HIPCC_FLAGS "${ROCMCC_FLAGS}")
else (${GPU_RUNTIME} MATCHES "HIP")
    set(HIPCC_FLAGS "${CUDACC_FLAGS}")
endif (${GPU_RUNTIME} MATCHES "HIP")

set_source_files_properties(${PENNANT_HIP_SRCS} PROPERTIES LANGUAGE ${GPU_RUNTIME})
set_source_files_properties(HydroGPU.hip PROPERTIES COMPILE_FLAGS ${HIPCC_FLAGS})

install(TARGETS pennant)

To test the cmake build system, do the following

mkdir build && cd build
cmake ..
make VERBOSE=1
ctest

Now testing for CUDA.

module load cuda

mkdir build && cd build
cmake -DCMAKE_GPU_RUNTIME=CUDA ..
make VERBOSE=1
ctest

For Perlmuter, make sure the right modules are loaded and it is very much the same.

module load PrgEnv-gnu/8.3.3
module load hip/5.4.3
module load PrgEnv-nvidia/8.3.3
module load cmake
export PATH=${PATH}:${HIP_PATH}

make

build/pennant test/sedovbig/sedovbig.pnt

make clean

```

```
mkdir build && cd build  
  
cmake -DCMAKE_GPU_RUNTIME=CUDA ..  
  
make VERBOSE=1  
  
./pennant ../test/sedovbig/sedovbig.pnt
```